Design Document

Totality AweSun

Bret Lorimore, Jacob Fenger, George Harder

February 14, 2017

CS 461 - Fall 2016

◆

Abstract

This document describes in detail the design for the various components of the North American Solar Eclipse 2017 senior capstone project. These components are described according to the IEEE 1016-2009 standard. There are three high level components detailed in this document, the Eclipse Image Processor, Eclipse Image Processor Manager, and Eclipse Simulator. The sections of this document are broken into subsections corresponding to these components as applicable.

David Konerding, Project Sponsor

Date

Bret Lorimore

Date

George Harder

Date

Jacob Fenger

Date

TABLE OF CONTENTS

1 DESIGN STAKEHOLDERS AND THEIR CONCERNS

The primary stakeholder in this project is David Konerding of Google. He is one of the managers of the Eclipse Megamovie project that is sponsoring this Senior Capstone project. David Konerding's concerns are listed below.

1.1 Image Processor

1.1.1

The image process needs to take in an image, identify if the image has a total solar eclipse, and if it does further process it so that it can be stitched into a timelapse movie by Eclipse Megamovie engineers.

1.1.2

The mean processing time for an image must be one second.

1.1.3

Images must be processed in no longer than five seconds.

1.1.4

Images must be filtered so that the processed images are only images of the eclipse at totality.

1.1.5

Only high quality images, defined as having a 50 pixel solar disk size and padding around the disk of 100 pixels, should be accepted by the image processor.

1.1.6

Once images have been filtered, they need to have metadata attached in a way that allows easy stitching of eclipse images.

1.1.7

The image processor needs to be able to be called by the image processor manager with appropriate input data.

1.1.8

Image processor needs to be able to use GPS EXIF information associated with images.

1.2   Image Processor Manager

1.2.1

Image processor manager should download images needing processing from Google Cloud Storage.

1.2.2

Image processor manager should invoke image processor with downloaded images.

1.2.3

Image processor manager should upload processed images and corresponding metadata - the output from the image processor - to Google Cloud Storage and Datastore, respectively.

1.2.4

Image processor should download/upload images at the same time the image processor is processing other images.

1.2.5

Image processor manager should invoke multiple image processor instances concurrently with different input images. The number of image processor instances launched should be determined by the number of cores on the host VM.

1.2.6

Image processor manager instances should be able to run alongside other image processor manager instances running on (potentially) different machines. These discrete instances should not attempt to process the same images.

1.3   Eclipse Simulator

1.3.1

The solar eclipse must be accurately simulated based on user entered location information.

1.3.2

Users will be able to view the eclipse via a draggable slider, clickable time step buttons, and a play functionality.

1.4

The user interface should be visually appealing, easy to navigate, and align with the theme of the larger project's site.

1.4.1

Simulator will only support locations within continental United States.

1.4.2

The simulator must load in less than 500ms given a 1-10 Mbps internet connection.

2 DESIGN VIEWPOINTS

2.1 Image Processor

2.1.1 Speed and Performance

Concerns: 1.1.2, 1.1.3
Elements: 4.1.1, 4.1.2, 4.1.3, 4.1.4
Analytical Methods: The primary criteria and methods in constructing this view is whether or not we are achieving the desired average speeds on our golden data set.
Viewpoint Source: George Harder

2.1.2 Accuracy

Concerns: 1.1.1, 1.1.4, 1.1.5, 1.1.8
Elements: 4.1.1, 4.1.2, 4.1.5, 4.1.6 4.1.9, 4.1.10
Analytical Methods: In constructing the corresponding view, we will be evaluating it based on whether or not the image processor is correctly identifying eclipses with at least 90% accuracy on our golden data set.
Viewpoint Source: George Harder

2.1.3 Input and Output

Concerns: 1.1.6, 1.1.7
Elements: 4.1.2, 4.1.4, 4.1.7, 4.1.8, 4.1.9
Analytical Methods: This view will be evaluating whether or not the input and output of the image process meet the specifications defined in the design document.
Viewpoint Source: George Harder

2.2   Image Processor Manager

2.2.1   Intra-instance Concurrency

Concerns: 1.2.4, 1.2.5

Elements: 4.2.1, 4.2.2, 4.2.3, 4.2.4, 4.2.5, 4.2.6, 4.2.7

Analytical Methods: Overall VM CPU utilization, overall VM network interface utilization. Both these values should be maximized as much as possible.

Viewpoint Source: Bret Lorimore

2.2.2   Inter-instance Concurrency and Synchronization

Concerns: 1.2.1, 1.2.3, 1.2.6

Elements: 4.2.1

Analytical Methods: No image should be successfully processed by multiple image processor instances, whether or not these run on the same VM or are managed by the same image processor manager.

Viewpoint Source: Bret Lorimore

2.2.3   Invocation

Concerns: 1.2.2

Elements: 4.2.6

Analytical Methods: Multiple image processor processes should be able to be easily launched concurrently with different input data.

Viewpoint Source: Bret Lorimore

2.3   Eclipse Simulator

2.3.1   Interface

Concerns: 1.3.1, 1.3.2

Elements: 4.3.1, 4.3.2, 4.3.4, 4.3.7

Analytical Methods: Interface should be appealing to the user as well as being responsive and fast. The look and feel of the user interface will be evaluated by the sponsor and deemed complete when it meets their desired expectations.

Viewpoint source: Jacob Fenger

2.3.2   Loading Performance

Concerns: 1.3.4

Elements: 4.3.3, 4.3.4, 4.3.5, 4.3.6, 4.3.7

Analytical Methods: The initial loading time of the simulator should be fast. Additionally, the interactions that the user has with the simulator should be responsive and should not show any significant slow downs.

Viewpoint source: Jacob Fenger

2.3.3 Simulation Accuracy

Concerns: 1.3.1, 1.3.3

Elements: 4.3.3, 4.3.5

Analytical Methods: In the simulator, the Sun and Moon display should reflect scientific accuracy when it comes to relative position and sizes. Additionally, the view of the Sun and Moon above the horizon shall be accurate.

Viewpoint source: Jacob Fenger

3 DESIGN VIEWS

3.1 Image Processor

3.1.1 Speed and Performance [Governed by Viewpoint 2.1.1]

The Eclipse Megamovie project expects to receive photographs on the order of hundreds of thousands from the numerous citizen photographers registered with the project. This being the case, it is of utmost importance that the image processor we are building for the project can process images in a timely manner. This is not only important to being able to build a movie from the images soon after the eclipse, but also because it prevents storage spaces on either end of the image processing pipeline from becoming clogged.

In order to achieve these design goals, we are designing a system that uses the OpenCV C++ library. This API provides us with high performing library methods, like Hough transforms, image crops, and image rotations that are necessary to the image processor's core functionality. In addition to the use of this language and library, we are also designing this system to be single threaded but to also interface with a manager that runs multiple instances of the image processor concurrently. This fact allows us to design the image process to process a single image quickly and accurately and leaves management of the application to a different art of the system.

3.1.2 Accuracy [Governed by Viewpoint 2.1.2]

From a high level, the most fundamental concern in the design of the image processor is that it can accurately identify an eclipse image. The other concerns associated with the design of the image processor are near meaningless if the application is not consistently identifying images that contain total eclipse.

To address this concern, and the related concerns around accepting only high quality images and determining the relative temporal and spatial positioning of images, we are designing this application with accuracy as

a primary goal. The meet this goal, the system will build upon an existing eclipse identification algorithm. This algorithm, with improvements we add ourselves, will be the basis for the parts of the system that we design to meet the accuracy requirements of the image processor.

3.1.3   Input and Output [Governed by Viewpoint 2.1.3]

The image processor is one component in a much larger system. For the system to function the image processor needs to interface with the other components in a well defined and seamless manner. In addition to speed and accuracy, we need to design the system with a view toward optimizing the way the image processor interacts with other components.

The ensure a smooth interface with the image processor manager we are designing the image processor to accept a well defined set of command line arguments including the path to the list of images to be processed, the path to an output directory, and the path to prefix image file names with that points to their location. The design of the image processor also takes into account the needs of the engineers handling the processed images. The image processor will write processed images and their metadata in a specific format to an output file.

3.2   Image Processor Manager

3.2.1   Maximal Utilization [Governed by viewpoint 2.2.1]

It is the goal of the image processor manager to maximize hardware utilization on the VMs on which it is running. This means that ideally, the VMs where the image processor manager and therefore the image processor are running will have an average of 100% CPU utilization on all cores at all times.The image processor application will be single threaded and thus by invoking multiple instances of this application, the image processor manager can increase utilization on multiple CPU cores.

The downloading of images to process and uploading of processed images are both very high latency operations. Therefore, instead of waiting for these downloads/uploads to complete with virtually zero CPU utilization in the meantime, the image processor manager can achieve greater average CPU utilization by completing these high latency tasks concurrently to processing other images.

3.2.2   Image Processing [Governed by viewpoint 2.2.2]

The ultimate goal of the image processor manager is to facilitate the processing of eclipse images by the image processor component of this project. This can be achieved by downloading images to process from the cloud, assembling them into a form that can be consumed by the image processor and then invoking that application with the downloaded images as input data.

3.2.3   Synchronization [Governed by viewpoint 2.2.3]

The image processor must process the images that are uploaded by users. These images are uploaded by the eclipsemega.movie website to Google Cloud Storage and metadata entries are also created for them in Google Cloud Datastore. In order for the image processor manager to invoke the image processor with these images, it must download them from Google Cloud.

The application to stitch processed images into movies, being developed by Google, expects to find processed eclipse images in Google Cloud Storage with corresponding metadata in Google Cloud Datastore. In order to meet this expectation, the image processor manager must upload the results of running the image processor to Google Cloud Storage and Google Cloud Datastore when ready.

To enable scalable image processing performance, it is desirable to enable many VMs to run image processor applications at once. In order to do this efficiently without redundancy, it is necessary to ensure that multiple VMs do not try to process the same images. For that reason, image processor manager instances must mark image files as pending processing before another instances of the image processor manager can queue them for processing. Without implementing this functionality, little to no performance improvements can be expected by deploying multiple image processor manager nodes.

3.3   Eclipse Simulator

3.3.1   User Interface View [Governed by Viewpoint 2.3.1]

The user interface shall utilize 2D animated depictions of the Sun and the Moon positioned based on the time the simulator is set to. In addition to the sun and moon, the user interface will contain stylized background imagery of a hillside. The simulator's controls will be split into a top bar and bottom bar. The top bar will have a location entry box, an expandable map, and a zoom button. The bottom bar will have a time slider, time up and down buttons, a play button, and a menu to select the play speed.

3.3.2   Operating Performance View [Governed by Viewpoint 2.3.2]

The simulator will have low loading times to ensure fast performance for most users. Additionally, the simulator will need to respond in a timely matter when users are interacting with the module.

3.3.3   Eclipse Accuracy View [Governed by Viewpoint 2.3.3]

The simulator shall be accurate enough for any location in the continental United States. This accuracy includes accurate relative Moon and Sun sizes, positions in the rendered scene, and positions relative to one another.

4 DESIGN ELEMENTS

4.1 Image Processor

4.1.1 OpenCV

Type: Library

Purpose: OpenCV is the computer vision library we will use to process images. This open source library can quickly crop and otherwise manipulate images. It best suits our needs for a computer vision utility.

4.1.2 C++

Type: Class

Purpose: This application will be written in C++ in order to give us better control over the speed at which the application runs and the necessary functionality for reading and writing files.

4.1.3 Image Processing Time

Type: Constraint

Purpose: This element exists because our requirements specify that the average time for an image to be processed must be less than one second.

4.1.4 Serial Image Processing

Type: Framework

Purpose: The image processor will process images one at a time because the Image Processor Manager handles all parallelization.

4.1.5 Hough Transform

Type: Procedure

Purpose: The Hough transform is an algorithm for identifying circles and lines in images. It will be used by the image processor to identify total eclipse images. OpenCV has the circular Hough transform built in.

4.1.6 Image Quality Error Checking

Type: Constraint

Purpose: The requirements of the image processor specify that only images with a solar disk of at least 50 pixels and 100 pixels of padding around the sun will be accepted.

4.1.7   Command Line Arguments

Type: Constraint

Purpose: The image processor needs to have a well defined set of command line arguments so that it can interface with the image processor manager without difficulty.

4.1.8   Data writer

Type: System

Purpose: This component is meant to encapsulate the functionality necessary to write the processed images and their associated metadata to an output file.

4.1.9   Image Data Structure

Type: Class

Purpose: This class will encapsulate the information and methods needed to manage the images we will be processing. It will also work closely with the Data Writer to write images and their metadata to files.

4.1.10   Solar Eclipse Image Standardisation and Sequencing (SEISS)

Type: Procedure

Purpose: The SEISS algorithm is an image processing algorithm that can identify images of eclipses, including eclipses at totality [1]. We will be basing part of the image processor off of this algorithm.

4.2   Image Processor Manager

4.2.1   Image List Downloader

Type: Subsystem

Purpose: The purpose of this subsystem is to download and return a list of images from Datastore to be processed by the image processor application. It will use Datastore transactions to ensure that each image in this list is marked as pending processing before it can be retrieved by another image processor manager instance. The max number of image files retrieved will be determined by the number of image processor processes that are going to be launched. This value will be a parameter of the image list downloader subsystem.

4.2.2   Image Downloader

Type: Subsystem

Purpose: The purpose of this subsystem is to download and save individual image files. It will accept the name of an image file to retrieve from Cloud Storage and a place to store this file, and will then download the file and save it to the desired location.

4.2.3 Image Download Manager

Type: System

Purpose: The purpose of this system is to coordinate design elements 4.2.1 and 4.2.2, the *Image List Downloader* and *Image Downloader*, respectively. It will use the *Image List Downloader* to retrieve a list of images to download and then will use the Python multiprocessing module to launch multiple instances of the *Image Downloader* subsystem concurrently to download the images in the list.

4.2.4 Result Uploader

Type: Subsystem

Purpose: The purpose of this subsystem is to upload an individual processed image to Google Cloud Storage and upload its metadata to Google Cloud Datastore.

4.2.5 Result Uploader Manager

Type: System

Purpose: The purpose of this system is to coordinate element 4.2.4, the *Results Uploader*. It will use the Python multiprocessing module to launch multiple instances of the *Results Uploader* concurrently to upload the output of the image processor.

4.2.6 Image Processor Invoker

Type: System

Purpose: The purpose of this system is to invoke multiple instances of the image processor application concurrently using the Python subprocess module. It will distribute images to process over multiple image processor processes to increase throughput. The number of image processor processes will be determined by the number of cores on the host VM.

4.2.7 Controller

Type: System

Purpose: The purpose of this system is to coordinate the three other systems that are part of the image processor manager, design elements 4.2.3, 4.2.5, and 4.2.6, the *Image Download Manager*, the *Result Uploader Manager*, and the *Image Processor Invoker*. It will run these systems in parallel so that images are downloaded/uploaded at the same time that other images are being processed.

4.3   Eclipse Simulator

4.3.1   Scalar Vector Graphics (SVG)

Type: System

Purpose: This element shall be used for the front-end display of the eclipse simulator. Two-dimensional images of the Sun and Moon will be altered based on how the user interacts with the module.

4.3.2   Cascading Style Sheets (CSS)

Type: System

Purpose: CSS helps with the front-end display of the simulator by helping produce better looking output.

4.3.3   Meuus JavaScript Library

Type: Library

Purpose: This library is used to compute the position of the Sun and Moon in the sky.

4.3.4   View

Type: Component

Purpose: Combined the HTML, SVG, and CSS elements for simulator display and interaction for the user.

4.3.5   Model

Type: Component

Purpose: Backend library in JavaScript used for computing eclipse information which will be passed to the controller. This entity utilized the Meuus JavaScript library for support.

4.3.6   Controller

Type: Component

Purpose: Controls the interaction between the model and view. Information will be passed between these entities.

4.3.7   Model-View-Controller Architecture

Type: Relationship

Purpose: This architecture is defined by the interactions of the model, view, and controller entities.

5 DESIGN OVERLAYS

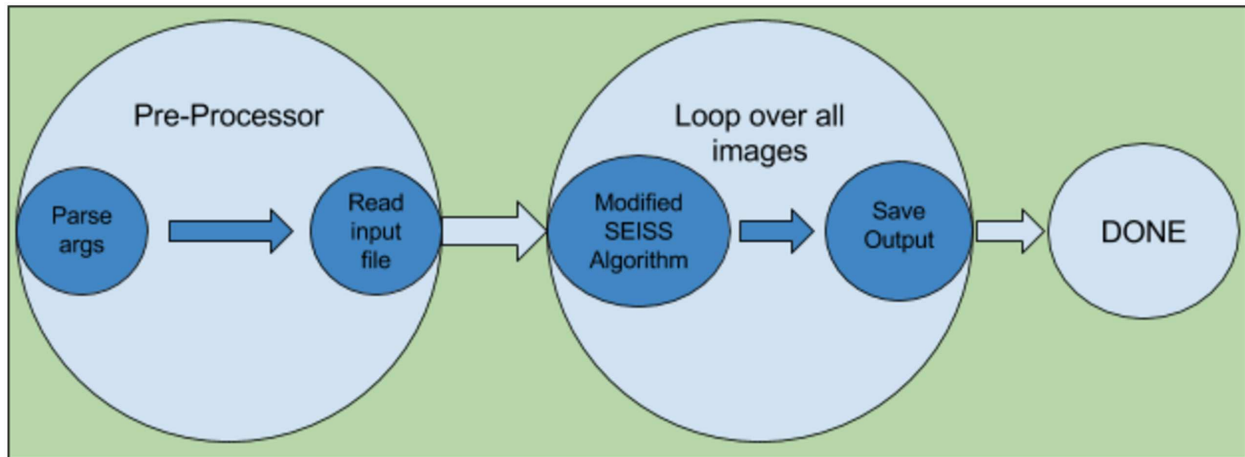## 5.1 Image Processor

Fig. 1. Image processor basic dataflow
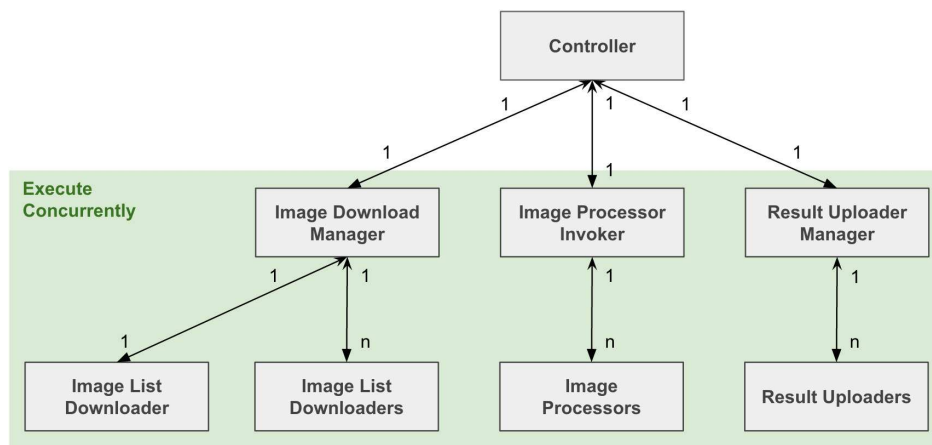
5.2 Image Processor Manager

Fig. 2. Image processor manager system diagram [Associated with Viewpoint 2.2.1]

5.3 Eclipse Simulator

Fig. 3. Eclipse simulator MVC architecture

6 DESIGN RATIONALE

6.1 Image Processor

The design of this system is based on needs surrounding accuracy, speed, and ability to interface with other components in the larger system. As was detailed in the technology review, we made certain decisions about what tools and algorithms to use based on the specific requirements of our system. The use of these tools has necessarily shaped the design of this application.

Additional information regarding design rationale can be found in the 'Purpose' section of the design elements.

6.2 Image Processor Manager

6.2.1 High Level System Design

From a high level the system is designed to be highly parallel with the aim of achieving as close to 100% CPU utilization on all cores as possible. In order to achieve this it is necessary to ensure that images to process are downloaded ahead of the time they're needed, this way the latency of their download time is hidden by processing other images. It is also necessary to ensure that processed images do not back up. The design elements outlined above fit together to form a system that naturally supports a very high level of concurrency. Subsystems to handle individual tasks are built and then multiple instances of these subsystems can be launched concurrently with different parameters. These subsystems are called by their parent systems.

6.2.2 Process Based Concurrency

As outlined in detail in our technology review, to achieve true multi-core concurrency in Python, use of process based concurrency such as that used in the built-in multiprocessing module is necessary. This is because multi-thread execution in Python is severely throttled by the Global Interpreter Lock (GIL) which essentially forces serialization of the execution of concurrent threads.

6.2.3   Motivation for separation of Image List Downloader and Image Downloader

Design element 4.2.1, the *Image List Downloader* is separated from design element 4.2.2, the *Image Downloader* as lists of all the images that need to be downloaded at a given time form relatively small collections of data. Therefore, it is more efficient to make a single Datastore query to retrieve all the images that need to be downloaded, rather than retrieving these one at a time using separate queries. As the image files themselves are quite large, these still need to be retrieved individually. In fact, the Google Cloud Client Library for Python does not support the download of multiple files from Cloud Storage using a single API call. On top of this constraint, downloading the actual image files individually ensures that we are able to establish a large number of concurrent connections to Cloud Storage to saturate the VM network interface as much as possible and achieve very high overall download speeds.

6.3   Eclipse Simulator

The goals for the simulator were to provide a fast and responsive experience to the user while providing scientifically accurate results. One main concern with this is that the time to compute information regarding the Sun and Moon must be quick enough to not pose any significant delay for the rest of the simulator.

We decided to utilize the MVC architecture since model and view components can be exchanged without compromising the whole system. System designer only need to account how the components interact with each other to update the system. In the technology review, we compared two libraries: Suncalc and Ephemeris. Initially, we chose SunCalc as the better library due to better documentation and ease of use. After further testing, we determined that the results of both of these libraries had unsatisfactory accuracy and opted for another choice: meuus.js. The Meuus library provides us with Sun/Moon positions that meet the requirements for accuracy of this simulator.

Additionally, we chose to utilize a scalar vector graphics format due to sub-element event processing and being easily to move the Sun and Moon around as animations.

7   DESIGN LANGUAGES

1)   UML Version 2.5

8 APPENDICES

8.1 Appendix I- Change History

1) November 30th, 2016- Document Created

8.2 Appendix II- Glossary of Terms

Eclipse Megamovie Project: The Eclipse Megamovie Project is a collaboration between Google and scientists from Berkeley and several other institutions with the aim of collecting large quantities of observations of the solar eclipse that will pass over the United States on August 21, 2017. The project will crowdsource photos of the eclipse from photographers at various points along the path of totality.

Metadata: In this document when we refer to metadata we are referring to any data associated with a file that is not that file itself - e.g. for a photo, the GPS coordinates at which that photo was taken are considered metadata.

Image processor instance/process: a single image processor process.

VM: A virtual machine - a hosted virtual server. In this document, when we refer to a VM we are specifically referring to a VM hosted in Google Cloud - this could be either a bare VM or a single node in a VM cluster.

Google Cloud Datastore: A fully managed, scheme-less NoSQL database solution offered from Google.

Google Cloud Storage: A fully managed file storage solution offered from Google, optimized for storing/accessing large files.

Mbps: Megabits per second. This is referring to download and upload speeds.

Ephemeris: An ephemeris is used to give positional information about astronomical objects. We are using a JavaScript library, meuus.js, as our ephemeris.

Machine: The term machine is used throughout this document in reference to a particular VM instance.

JPEG/JPG: JPEG is a lossy compression technique for images. When we refer to JPEG/JPG files in this document we are referring to image files compressed in this method with the .jpeg or .jpg file extension.

PNG: PNG refers to the Portable Network Graphics image file format. Images in the PNG format are frequently referred to as "PNGs" and are saved with the .png file extension.

EXIF: EXIF refers to the Exchangeable Image File Format, a standard media file format. Within the scope of this document EXIF generally refers to metadata associated with images processed by the image processor.

SEISS: Solar Eclipse Image Standardisation and Sequencing. This refers to an algorithm developed by Krista et al. [1] that seeks to identify eclipse images and classify them by the phase that the eclipse is in. We will be using and modifying this algorithm to suit the needs of this project.

REFERENCES

[1] K. Larisza and S. McIntosh, "The standardisation and sequencing of solar eclipse images for the eclipse megamovie project."