

Requirements

Totality AweSun

Bret Lorimore, Jacob Fenger, George Harder

November 4th, 2016

CS 461 - Fall 2016

◆

Abstract

On August 21, 2017 a total solar eclipse will pass over the United States. The path of totality will stretch from Oregon to South Carolina. There has not been a total solar eclipse like this, crossing the country from coast to coast, since the eclipse of 1918. The Eclipse Megamovie Project is a collaboration between Google and scientists from UC Berkeley and several other institutions with the aim of compiling a large dataset of eclipse observations. Acquiring coronal data is of particular interest as the corona is not normally visible from Earth. Specifically, the project will crowdsource photos of the eclipse from photographers at various locations along the path of totality. These images will be aligned spatially and temporally and stitched into a unique movie that shows the eclipse over a period of 1.5 hours as it passes across the United States. Additionally, the complete photo dataset will be open sourced so that independent researchers may do their own analysis.

Google will contribute applications providing, among other things, backend image processing, photo upload capabilities, and static informational content. This senior capstone project will consist of three distinct sub-projects, specifically, implementing/modifying an image processing algorithm facilitating the classification and alignment of solar eclipse images before they are stitched into a movie, a manager to run this image processing app, and a location-based eclipse simulator.

David Konderding, Project Sponsor

Date

Bret Lorimore

Date

George Harder

Date

Jacob Fenger

Date

TABLE OF CONTENTS

1	Introduction	3
1.1	Purpose	3
1.2	Scope	3
1.3	Definitions, acronyms, and abbreviations	3
1.4	References	4
1.5	Overview	4
2	Overall Description	5
2.1	Product perspective	5
2.2	Product functions	5
2.3	User characteristics	6
2.4	Constraints	6
2.5	Assumptions and dependencies	6
2.6	Apportioning of requirements	6
3	Specific requirements	7
3.1	External Interfaces	7
3.1.1	Eclipse Simulator	7
3.1.2	Eclipse Image Processor	7
3.1.3	Eclipse Image Processor Manager	8
3.2	Functional Requirements	8
3.2.1	Eclipse Simulator	8
3.2.2	Eclipse Image Processor	9
3.2.3	Eclipse Image Processor Manager	10
3.3	Performance Requirements	10
3.3.1	Eclipse Simulator	10
3.3.2	Eclipse Image Processor	10
3.3.3	Eclipse Image Processor Manager	10
4	Supporting Information	12
4.1	Appendix	12

1 INTRODUCTION

1.1 Purpose

The purpose of this software requirements specification (SRS) is to describe in detail the Eclipse Image Processor, Eclipse Image Processor Manager, and the Eclipse Simulator that our group will produce. By writing these requirements down and agreeing to them with our sponsor both parties will have a clear understanding of what the finished product will be and what it will be able to do. The intended audience for this SRS is our sponsor, the Senior Capstone Instruction Team, and ourselves.

1.2 Scope

We are producing ~~two~~ three products: an Eclipse Image Processor, Eclipse Image Processor Manager, and an Eclipse Simulator. The Eclipse Image Processor will ingest images ~~find the Sun and the Moon in the images, crop the image around these bodies, and extract information about the relative position of the Sun and the Moon in these images~~ and if they are of a total solar eclipse, align/crop these images consistently. Images of non-total solar eclipses will marked as such and otherwise ignored. The Eclipse Image Processor Manager will download images to be processed by the Image Processor, run the Image Processor on these images, and collect the results of running the Image Processor and upload them to Google Cloud. The Eclipse Simulator will provide users with a 2D visual representation of the eclipse from a specified location from 12 hours before it occurs to 12 hours after.

1.3 Definitions, acronyms, and abbreviations

Eclipse Megamovie Project: The Eclipse Megamovie Project is a collaboration between Google and scientists from Berkeley and several other institutions with the aim of collecting large quantities of observations of the solar eclipse that will pass over the United States on August 21, 2017. The project will crowdsource photos of the eclipse from photographers at various points along the path of totality.

EXIF: EXIF refers to the Exchangeable Image File Format, a standard media file format. Ancillary data tags associated with other media files are frequently referred to as "EXIF fields", "EXIF data", etc. This is how "EXIF" will be used within the scope of this document.

GPS: GPS stands for global positioning system. We will refer to the latitude and longitude information obtained from a global positioning system as "GPS coordinates", "GPS data", etc. throughout this document.

JPEG/JPG: JPEG is a lossy compression technique for images. When we refer to JPEG/JPG files in this document we are referring to image files compressed in this method with the .jpeg or .jpg file extension.

PNG: PNG refers to the Portable Network Graphics image file format. Images in the PNG format are frequently referred to as "PNGs" and are saved with the .png file extension.

1.4 References

This SRS makes reference to a report by Larisza D. Krista and Scott W. McIntosh titled "The Standardisation and Sequencing of Solar Eclipse Images for the Eclipse Megamovie Project." This technical report was produced as a collaboration between scientists at the University of Colorado at Boulder, the National Center for Atmospheric Research and the National Oceanic and Atmospheric Administration. This paper can be found on [arxiv.org](https://arxiv.org/abs/1508.01713).

1.5 Overview

The remainder of this SRS contains an overall description of the Eclipse Image Processor, Eclipse Image Processor Manager, and Eclipse Simulator systems in section 2. Following these descriptions are specific requirements for the systems in section 3.

2 OVERALL DESCRIPTION

2.1 Product perspective

- 1) These products, the Eclipse Image Processor, Eclipse Image Processor Manager, and the Eclipse Simulator, are all components of the larger Eclipse Megamovie Project. ~~They are designed to operate as wholly independent modules that can be "plugged into" the existing Eclipse Megamovie codebase.~~ The Eclipse Image Processor will be a binary that receives images from another application - the Eclipse Image Processor Manager - processes them, and exports them. The Eclipse Image Processor and its manager will work together and form a standalone module that can be deployed to multiple Virtual Machines as a Docker container. The Eclipse Simulator will be a standalone JavaScript module that can be added to an existing webpage.
- 2) The Eclipse Image Processor and Image Processor Manager do not directly interface with the user, they are backend systems that will be running asynchronously in Google Cloud. ~~It resides behind the front end of the Eclipse Megamovie website.~~ The Eclipse Simulator does interface directly with the user. It should function on most modern internet browsers (Chrome, Firefox, Safari). The simulator will appear to the user as a 2D animated depiction of the Sun and the Moon as they appear at the specified time and location. The simulator will also have background imagery in addition to the Sun and the Moon. Besides the images, the simulator will have a time slider, a location input, and a time display.
- 3) This system does not interface with hardware.
- 4) The Eclipse Image Processor and Image Processor Manager will be designed to run on Ubuntu 16.04. It is necessary for these applications to be compatible with this operating system because the machines that will be running the them also run Ubuntu 16.04. The Eclipse Simulator is a JavaScript module that will work on modern browsers like Chrome, Firefox and Safari. We expect our users will use these popular browsers so it is necessary for our product to interface with them.

2.2 Product functions

- 1) The Eclipse Simulator will be a standalone JavaScript module enabling users to "preview" the eclipse. It will be designed in a stylized, 2D manner. The simulator will incorporate a time slider that allows users to simulate the eclipse in a time window spanning from 12 hours before the eclipse to 12 hours after it. As users drag the time slider, the eclipse will animate in the simulator window. The view of the eclipse which users are presented will be specific to the selected location.
- 2) The Eclipse Image Processor application will ingest photos of the eclipse, determine if they are images of a total solar eclipse, and if so, align them spatially ~~and categorize them to help recover temporal ordering.~~ These aligned ~~and categorized~~ images will be cropped so that the Sun occupies the same amount of space in each image and will be exported. All data that is required to take the raw input images and produce the exported images will be saved to a data file. Additionally, select EXIF information will be extracted from the image files and included in this output data file. Images that are deemed as being of "non-total" solar eclipses will be marked as such and otherwise ignored.

- 3) The Eclipse Image Processor Manager will be responsible for managing the Eclipse Image Processor application. This includes collecting images from Google Cloud for the Image Processor to process, invoking the Image Processor with these images as input, and collecting the output of the Image Processor and uploading it to Google Cloud. The Image Processor and Image Processor Manager will be deployed together in a single docker container to Google Container Engine Clusters (of VMs). An important role of the Image Processor Manager is that it will be responsible for ensuring that the compute resources on the host VMs are as saturated as possible. This means invoking multiple Image Processor processes concurrently, while at the same time downloading the next images to be processed and uploading the already processed images.

2.3 User characteristics

- 1) The Eclipse Image Processor application will be used by ~~Google Engineers~~ the members of this project.
- 2) The Eclipse Image Processor Manager application will be used by Google Engineers.
- 3) The Eclipse Simulator application will be used by the general public. No unusual technical/scientific knowledge is expected of these users. It is assumed however, that these users are familiar with the internet and web browsers.

2.4 Constraints

None.

2.5 Assumptions and dependencies

- 1) This SRS assumes the availability of Ubuntu 16.04.
- 2) This SRS assumes the availability of Google Cloud Platform n1-standard-4 virtual machines.
- 3) This SRS assumes the ability to interface with Google Cloud from Python, for example with the Google Cloud Client Library for Python.
- 4) This SRS assumes the availability of the OpenCV computer vision library.

2.6 Apportioning of requirements

See Gantt Chart in Appendix.

3 SPECIFIC REQUIREMENTS

3.1 External Interfaces

3.1.1 Eclipse Simulator

- 1) The simulator is a standalone JavaScript module that can be included on an existing webpage.
- 2) Users can select the location from which to simulate the eclipse. This can be entered at any point while using the simulator.
 - a) Location can be entered as: latitude/longitude, address, zip code, city name, state name.
 - b) Initial simulator location can be programmatically set as initialization parameter.
- 3) Users will be able to adjust the simulator time from 12 before the eclipse to 12 hours after it.
 - a) Time can be advanced via a draggable slider or clickable buttons.

3.1.2 Eclipse Image Processor

- 1) The Image Processor will be compatible with Ubuntu 16.04 and will include a script that to install all dependencies/build the binary.
- 2) The application will accept the following input as command line arguments:
 - a) Required: image_list_file
 - i) Absolute or relative (to the directory the binary was invoked from) path to file containing a list of image filenames with no directory prefix.
 - b) Required: output_dir
 - i) Directory to write output files to.
 - c) Optional: image_path_prefix
 - i) Absolute or relative (to the directory the binary was invoked from) path to prepend to each image filename in image_file_list. Defaults to "/".
- 3) The application will accept JPEG (.jpeg/.jpg) and PNG (.png) image files.
 - a) Images of an invalid format will be disregarded and an error message will be written to stderr.
- 4) The application will write the following output to the output_dir directory:
 - a) image_transformations.txt
 - i) File containing one line per image processed with the following values (comma separated):
 - A) processed_image: processed image filename
 - B) ~~image_type: image type (FULL_DISK/TOTALITY/etc.); see requirement #2 of 3.2.2~~
 - C) ~~rot_angle: angle original image was rotated (degrees)~~
 - D) crop_topl_x: x coordinate of top left corner of cropped image (refers to rotated image)
 - E) crop_topl_y: y coordinate of top left corner of cropped image (refers to rotated image)
 - F) crop_botr_x: x coordinate of bottom right corner of cropped image (refers to rotated image)
 - G) crop_botr_y: y coordinate of bottom right corner of cropped image (refers to rotated image)

- H) ~~rel_center_offset: relative offset of solar/lunar disk centers, see requirement #3 of 3.2.2 (optional, only included for CRESCENT type images)~~
 - I) ~~diamond_rel_size: size of diamond relative to the size of the solar disk, see requirement #4 of 3.2.2 (optional, only for DIAMOND_RING type images)~~
 - J) ~~timestamp: timestamp at which image was taken (optional, only included when images have an EXIF timestamp field)~~
 - K) ~~eclipse_path_percent: percentage through eclipse totality path at which image was taken, see requirement #8 of 3.2.2 (optional, only included for images with GPS coordinate EXIF field)~~
 - L) ~~rejected_reason: reason a particular image was rejected, see requirement #10 of 3.2.2.~~
- b) Processed image files
- i) Processed image files will be saved into a sub-directory of output_dir called "images" and will be named as follows: *_pp.[png|jpeg|jpg].
- 5) Images that are rejected will ~~not~~ be added to the image_transformations.txt file but will not be cropped/exported to the images directory. ~~When an image is rejected a message with this information will be printed to stdout.~~
- 6) The application will format all log messages as follows:
- a) img_preproc:level:timestamp:message
 - i) Values of level: ERROR/WARNING/INFO/DEBUG
 - ii) Value of timestamp: current timestamp
 - iii) Value of message: specific logging message

3.1.3 Eclipse Image Processor Manager

- 1) The Eclipse Image Processor Manager will collect images to be processed from Google Cloud Storage. It will obtain lists of files to download and their respective metadata from Google Cloud Datastore. It will be aware of the scheme of the Datastore NoSQL database in order to read/write image metadata.
- 2) The Eclipse Image Processor Manager will invoke instances of the Eclipse Image Processor and communicate with them via a command-line interface as well as data files. The format of these data files, as well as the command-line interface is described in 3.1.2.

3.2 Functional Requirements

3.2.1 Eclipse Simulator

- 1) Displayed solar/lunar placement will be based on location and time and will account for edge cases like when the location is not in the path of totality. ~~For example, if the location is on the opposite side of the world as the eclipse, the simulator should shift to a night time display.~~
- 2) The simulator location will be restricted to North America as this is where the eclipse will be visible from.
- 3) The Simulator will display a timestamp indicating the time that the simulator is set to. ~~Simulator will display the local time that the simulator is set to, e.g. there is a well defined time associated with the user selecting Corvallis, Oregon as their location and a simulator time of 3:13 (3 hours 13 minutes) before the eclipse. This time should be displayed on the simulator.~~

3.2.2 Eclipse Image Processor

- 1) Invalid JPEG and PNG files (e.g. cannot be opened by OpenCV, width/height equal to 0px, etc.) will be ignored and an error message will be written to stderr.
- 2) ~~The application will classify the input images as being one of the following types:~~
 - a) ~~FULL_DISK~~
 - i) ~~Image of an unobscured solar disk.~~
 - b) ~~TOTALITY~~
 - i) ~~Image of a total solar eclipse.~~
 - c) ~~CRESCENT~~
 - i) ~~Image of a partially eclipsed Sun, creating a "crescent" shape.~~
 - d) ~~DIAMOND_RING~~
 - i) ~~Image of a nearly fully eclipsed Sun where there is one "hot spot" on the Sun's perimeter. This hot spot along with the Sun's perimeter have the shape of a diamond ring.~~
- 3) ~~For images of type CRESCENT, the application will compute/export a delta of the position of the center of the Sun and the Moon relative to the size of the solar disk. This delta will be a signed value based on the Sun's position, i.e. if the Moon is to the left of the Sun (in the cropped/rotated image) the delta will be negative and conversely, if the Moon is to the right of the Sun the delta value will be positive.~~
- 4) ~~For images of type DIAMOND_RING, the application will compute/export the size of the "diamond" relative to the size of the solar disk.~~
- 5) Images where the solar disk has a radius of less than 50px will be rejected (see requirement #5 in 3.1.2).
- 6) The application will crop the images to be square with the Sun centered. The images will be cropped so that there is a 100px pad between the solar perimeter and the edge of the image on all sides.
 - a) Images that do not have enough room around the Sun to allow for the pad described above will be rejected, see requirement #5 in 3.1.2.
- 7) ~~The application will rotate DIAMOND_RING and CRESCENT type images so that they are aligned horizontally, as described by Krista et al.~~
- 8) ~~For images with GPS EXIF information,~~ Based on images' GPS EXIF information, the application will compute/export the percentage through the eclipse's path of totality at which the image was taken. 0% will be defined as the westmost point on that path of totality that is over land, this point is on the west coast of Oregon. 100% will be defined as the eastmost point on the path of totality that is over land, this point is on the east coast of South Carolina. The application will compute the point on the path of totality nearest the point where the image was taken. This point will be used to compute the percentage through the path of totality at which the image was taken.
- 9) ~~Images with GPS EXIF information that are~~ that were not taken on the path of totality will be rejected, see requirement #5 in 3.1.2.
- 10) Images can be rejected by the Image Processor. When an image is rejected, the reason it was rejected will be

noted and exported, see requirement #5 in 3.1.2. The possible reasons for which an image can be rejected will include at least the following:

- a) Invalid image file, e.g. image file is corrupted and cannot be opened.
- b) Solar disk radius is too small.
- c) There is not enough padding around the solar disk.
- d) Image was not taken on the path of totality.

3.2.3 Eclipse Image Processor Manager

- 1) The application will be designed so that multiple instances of it can run concurrently with no communication between them. These distinct managers will not attempt to process the same images.
- 2) The application will record the time at which it "checks-out" image files for processing in Datastore. This enables future processing of images that were downloaded by an Image Processor Manager instance that later became delinquent.
- 3) The application will group files/metadata downloaded from Google Cloud into directories corresponding to individual workloads that will be delegated to distinct Image Processor instances.
- 4) When an Image Processor instance completes its workload, its output data will be collected by its parent Image Processor Manager and uploaded to Google Cloud. This will include processed image files and metadata that is saved to Datastore.
- 5) The application will kill delinquent Image Processor Instances and upload the results of processing any images that were processed successfully. Images that were not processed successfully will be marked as such in Datastore.
- 6) The application will be highly parallel. Multiple Eclipse Image Processor application instances will be launched concurrently, and while these are working, the next images to be processed will be downloaded and previous results will be uploaded.
 - a) The number of Image Processor instances launched will be determined by the number of cores on the host VM. *Note: this does not mean that given an n core machine, n Image Processor instances will be launched - some cores must be reserved for downloading/preparing input data, and processing/uploading results.*

3.3 Performance Requirements

3.3.1 Eclipse Simulator

- 1) All simulator resources will load in less than 500ms given a 1-10 Mbps internet connection.

3.3.2 Eclipse Image Processor

- 1) The application should take less than 5 seconds to process an image when running on a Google Cloud Platform n1-standard-4 virtual machine.

3.3.3 Eclipse Image Processor Manager

- 1) The Image Processor Manager will saturate either the CPUs or network interface of the host VM to which it is deployed.

- The reason for the "or" in the above statement is that either the time to process images, or the time to download them and upload their processed counterparts will be a bottleneck depending on the VM configuration. The goal here is to get as close to network saturation and 100% utilization of all CPU cores as possible. If all CPU cores are running at nearly 100% but the network interface is not saturated, then VMs with more cores should be used. Conversely, if the network interface is saturated but not all cores are running at 100%, VMs with fewer cores should be used.

